

## Importing Important Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

## Load the dataset

```
In [2]: df=pd.read_csv('thyroid_disease.csv')
df
```

Out[2]:

	age	sex	on_thyroxine	query_on_thyroxine	on_antithyroid_meds	sick	pregnant	thyroid_surgery	l131_treatment	query_hypothyroid	...	TT4	T4U_me
0	29	F	f	f	f	f	f	f	f	f	t ...	NaN	
1	29	F	f	f	f	f	f	f	f	f	f ...	128.0	
2	41	F	f	f	f	f	f	f	f	f	f ...	NaN	
3	36	F	f	f	f	f	f	f	f	f	f ...	NaN	
4	32	F	f	f	f	f	f	f	f	f	f ...	NaN	
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9167	56	M	f	f	f	f	f	f	f	f	f ...	64.0	
9168	22	M	f	f	f	f	f	f	f	f	f ...	91.0	
9169	69	M	f	f	f	f	f	f	f	f	f ...	113.0	
9170	47	F	f	f	f	f	f	f	f	f	f ...	75.0	
9171	31	M	f	f	f	f	f	f	f	f	t ...	66.0	

9172 rows × 31 columns

## dropping redundant attributes from thyroidDF dataset

```
In [3]: df.drop(['TSH_measured', 'T3_measured', 'TT4_measured', 'T4U_measured', 'FTI_measured', 'TBG_measured',
               'patient_id', 'referral_source'], axis=1, inplace=True)
```

In [4]: df

Out[4]:

	age	sex	on_thyroxine	query_on_thyroxine	on_antithyroid_meds	sick	pregnant	thyroid_surgery	l131_treatment	query_hypothyroid	...	tumor	hypopiti
0	29	F	f	f	f	f	f	f	f	f	t ...	f	
1	29	F	f	f	f	f	f	f	f	f	f ...	f	
2	41	F	f	f	f	f	f	f	f	f	f ...	f	
3	36	F	f	f	f	f	f	f	f	f	f ...	f	
4	32	F	f	f	f	f	f	f	f	f	f ...	f	
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9167	56	M	f	f	f	f	f	f	f	f	f ...	f	
9168	22	M	f	f	f	f	f	f	f	f	f ...	f	
9169	69	M	f	f	f	f	f	f	f	f	f ...	f	
9170	47	F	f	f	f	f	f	f	f	f	f ...	f	
9171	31	M	f	f	f	f	f	f	f	f	t ...	f	

9172 rows × 23 columns

## Checking Null Values in Data set

```
In [5]: df.isnull().sum()
```

```
Out[5]: age                0
sex                307
on_thyroxine       0
query_on_thyroxine 0
on_antithyroid_meds 0
sick               0
pregnant          0
thyroid_surgery    0
I131_treatment     0
query_hypothyroid  0
query_hyperthyroid 0
lithium            0
goitre            0
tumor             0
hypopituitary     0
psych             0
TSH               842
T3               2604
TT4              442
T4U              809
FTI              802
TBG              8823
target            0
dtype: int64
```

```
In [6]: df.shape
```

```
Out[6]: (9172, 23)
```

```
In [7]: df.columns
```

```
Out[7]: Index(['age', 'sex', 'on_thyroxine', 'query_on_thyroxine',
              'on_antithyroid_meds', 'sick', 'pregnant', 'thyroid_surgery',
              'I131_treatment', 'query_hypothyroid', 'query_hyperthyroid', 'lithium',
              'goitre', 'tumor', 'hypopituitary', 'psych', 'TSH', 'T3', 'TT4', 'T4U',
              'FTI', 'TBG', 'target'],
              dtype='object')
```

```
In [8]: df['target'].unique()
```

```
Out[8]: array(['-', 'S', 'F', 'AK', 'R', 'I', 'M', 'N', 'G', 'K', 'A', 'KJ', 'L',
              'MK', 'Q', 'J', 'C|I', 'O', 'LJ', 'H|K', 'D', 'GK', 'MI', 'P',
              'FK', 'B', 'GI', 'C', 'GKJ', 'OI', 'D|R', 'E'], dtype=object)
```

## re-mapping target values to diagnostic groups

```
In [9]:
```

```
diagnoses = {'-': 'negative',
             'A': 'hyperthyroid',
             'B': 'hyperthyroid',
             'C': 'hyperthyroid',
             'D': 'hyperthyroid',
             'E': 'hypothyroid',
             'F': 'hypothyroid',
             'G': 'hypothyroid',
             'H': 'hypothyroid'}

df['target'] = df['target'].map(diagnoses) # re-mapping
# dropping observations with 'target' null after re-mapping
df.dropna(subset=['target'], inplace=True)
```

dataset initial summary

In [10]:

df.info()

<class 'pandas.core.frame.DataFrame'>  
Int64Index: 7546 entries, 0 to 9171  
Data columns (total 23 columns):  
# Column Non-Null Count Dtype  
--- --- -  
0 age 7546 non-null int64  
1 sex 7296 non-null object  
2 on\_thyroxine 7546 non-null object  
3 query\_on\_thyroxine 7546 non-null object  
4 on\_antithyroid\_meds 7546 non-null object  
5 sick 7546 non-null object  
6 pregnant 7546 non-null object  
7 thyroid\_surgery 7546 non-null object  
8 l131\_treatment 7546 non-null object  
9 query\_hypothyroid 7546 non-null object  
10 query\_hyperthyroid 7546 non-null object  
11 lithium 7546 non-null object  
12 goitre 7546 non-null object  
13 tumor 7546 non-null object  
14 hypopituitary 7546 non-null object  
15 psych 7546 non-null object  
16 TSH 6824 non-null float64  
17 T3 5337 non-null float64  
18 TT4 7192 non-null float64  
19 T4U 6870 non-null float64  
20 FTI 6877 non-null float64  
21 TBG 259 non-null float64  
22 target 7546 non-null object  
dtypes: float64(6), int64(1), object(16)  
memory usage: 1.4+ MB

distributions of numeric variables

In [11]:

df.describe()

Out[11]:

	age	TSH	T3	TT4	T4U	FTI	TBG
count	7546.000000	6824.000000	5337.000000	7192.000000	6870.000000	6877.000000	259.000000
mean	78.013782	5.421753	2.020935	105.203373	0.967322	110.571745	22.955019
std	1305.258137	26.080471	0.809865	32.606462	0.162315	36.600867	6.088392
min	1.000000	0.005000	0.050000	2.000000	0.190000	1.400000	0.100000
25%	37.000000	0.570000	1.600000	87.000000	0.870000	93.000000	20.000000
50%	55.000000	1.400000	2.000000	103.000000	0.960000	108.000000	23.000000
75%	67.000000	2.700000	2.300000	121.000000	1.060000	125.000000	27.000000
max	65526.000000	530.000000	18.000000	430.000000	2.120000	839.000000	45.000000

Observations:

--> Max age value is 65,526 years old. Either that person is really really old or this is a mistake. There are likely more inconsistencies like this one throughout the data. --> Persons with age above 100 will be removed since they have target negative and we wont be losing to much information by omitting them

inspecting observations with age > 100

In [12]:

df[df.age > 100]

Out[12]:

	age	sex	on_thyroxine	query_on_thyroxine	on_antithyroid_meds	sick	pregnant	thyroid_surgery	l131_treatment	query_hypothyroid	...	tumor	hypop
2976	455	F	f		f	f	f	f	f	f	...	f	
5710	65511	M	f		f	f	f	f	f	f	...	f	
6392	65512	M	f		f	f	f	f	f	f	...	f	
8105	65526	F	f		f	f	f	f	f	f	...	f	

4 rows × 23 columns

**changing age of observations with ('age' > 100) to null**

```
In [13]: df['age'] = np.where((df.age > 100), np.nan, df.age)
```

```
In [14]: df['target'].value_counts()
```

```
Out[14]: negative      6771  
hypothyroid    593  
hyperthyroid   182  
Name: target, dtype: int64
```

**Exploratory Data Analysis**

```
In [15]: # ---> We begin our EDA by looking at the distribution of Hormone Levels in blood for each of our target classes.  
# ---> This helps us get an idea for how good of a predictor each of these attributes can be.
```

## setting up grid for multiple seaborn plots

In [16]:

```
fig, axes = plt.subplots(3,2,figsize=(20,16))
fig.suptitle('Numerical Attributes vs. Target')
sns.set_style('whitegrid');

# TSH vs. 'target'
sns.stripplot(x=df.target, y=df.TSH, linewidth=0.6, jitter= 0.3, ax=axes[0, 0])

# T3 vs. 'target'
sns.stripplot(x=df.target, y=df.T3, linewidth=0.6, jitter= 0.3, ax=axes[0, 1])

# TT4 vs. 'target'
sns.stripplot(x=df.target, y=df.TT4, linewidth=0.6, jitter= 0.3, ax=axes[1, 0])

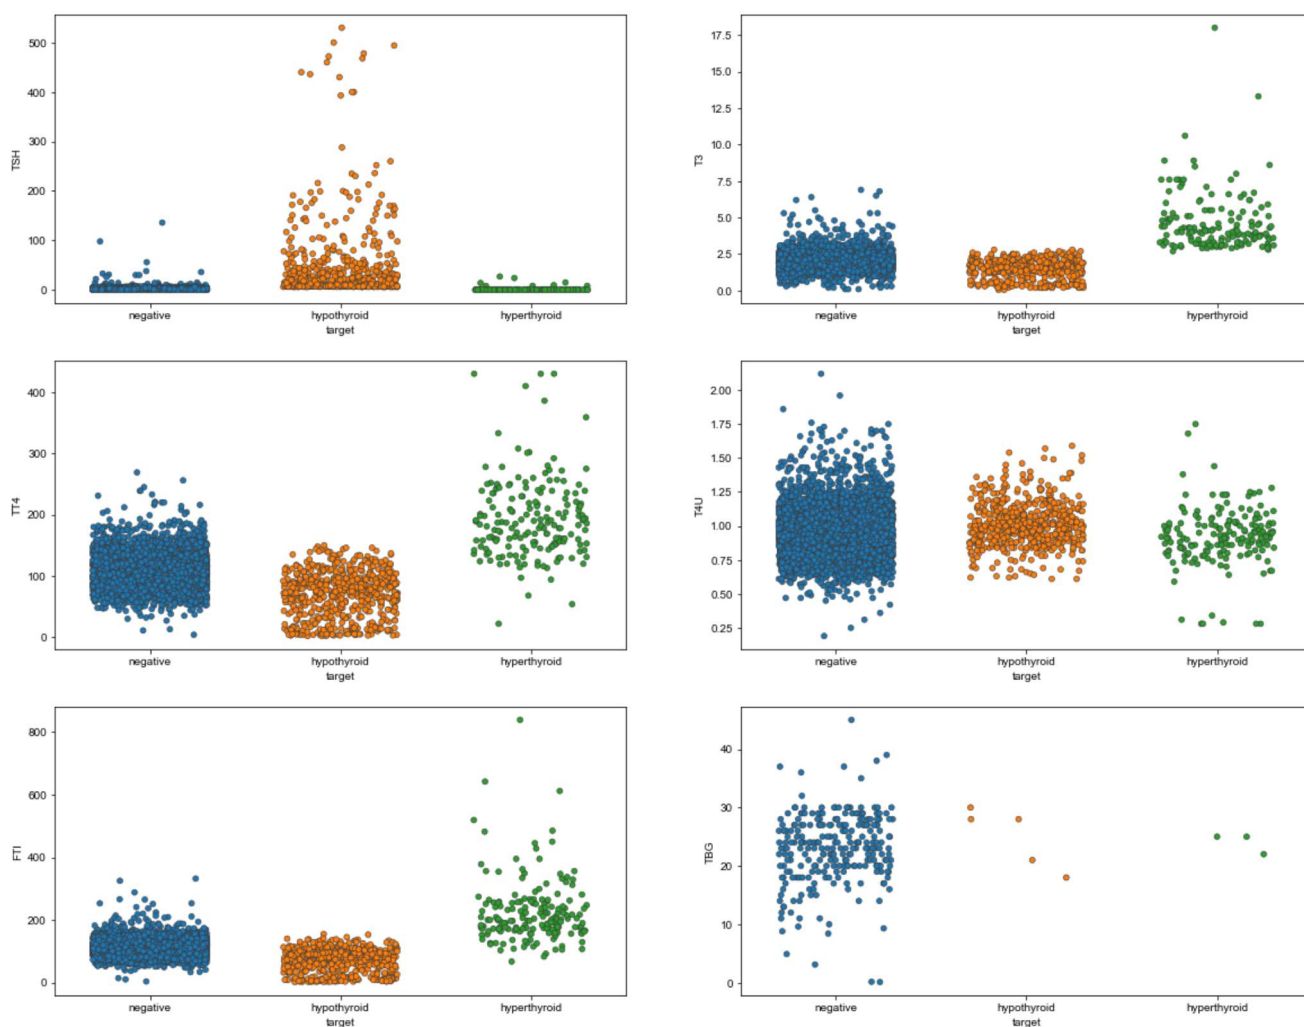
# T4U vs. 'target'
sns.stripplot(x=df.target, y=df.T4U, linewidth=0.6, jitter= 0.3, ax=axes[1, 1])

# FTI vs. 'target'
sns.stripplot(x=df.target, y=df.FTI, linewidth=0.6, jitter= 0.3, ax=axes[2, 0])

# TBG vs. 'target'
sns.stripplot(x=df.target, y=df.TBG, linewidth=0.6, jitter= 0.3, ax=axes[2, 1])
```

Out[16]: &lt;AxesSubplot:xlabel='target', ylabel='TBG'&gt;

Numerical Attributes vs. Target



## Observations:

In [17]:

```
# --> Immediately we can hypothesize that FTI, T3, and TT4 will be good feature additions to our models.
# TSH looks like it might be good as well but we need to handle the outliers for 'target' hypo and analyze the attribute dis
# --> This is all in-line with the knowledge discovered about Hormone Level tests during our initial research
```

Let's continue by creating a pairplot of our numeric variables and seeing if we can spot any clusters forming between variables.

In [18]:

```
numericalDF = df[['age', 'TSH', 'T3', 'TT4', 'T4U', 'FTI', 'target']].copy()
sns.set_style('whitegrid');
sns.pairplot(numericalDF, hue='target', height=3);
plt.show()
```



In the diagonals of the pairplot we can see the distributions of each numeric variable with respect to one another. It is apparent how unbalanced the dataset is, with so many negative 'target' compared to hypothyroid or hyperthyroid.

## Observations:

We can see that for some Hormone test vs others there are nice clusters that form. This is encouraging because it means that they do a good job at separating out each of our target classes. FTI vs T3 FTI vs T4U FTI vs age T4U vs TT4 TT4 vs age TT4 vs T3 There is severe target class imbalance... but we knew this from the start. It is normal for this type of data. We will have to treat with resampling protocol as well as using models that handle this well.

## Investigating Feature Correlations


```
In [19]: # Now Lets give some attributes and take a look at the correlation
# between of our numerical attributes between one another.
```

```
In [20]: df_mat = df.corr()
round(df_mat,2)
```

Out[20]:

	age	TSH	T3	TT4	T4U	FTI	TBG
age	1.00	-0.02	-0.16	-0.03	-0.09	0.02	0.10
TSH	-0.02	1.00	-0.19	-0.32	0.10	-0.32	-0.12
T3	-0.16	-0.19	1.00	0.58	0.17	0.52	-0.01
TT4	-0.03	-0.32	0.58	1.00	0.30	0.79	0.01
T4U	-0.09	0.10	0.17	0.30	1.00	-0.26	0.57
FTI	0.02	-0.32	0.52	0.79	-0.26	1.00	-0.15
TBG	0.10	-0.12	-0.01	0.01	0.57	-0.15	1.00

```
In [21]: sns.heatmap(df_mat);
```



## Handling Inconsistencies

### Investigating Outliers

The code below calculates the Inter-quartile ranges for our Hormone test numeric attributes in order to determine mild and severe outliers. Depending on the severity of the issue we will decide how to handle them in the next section.

```

In [22]: # TSH
Q1_TSH = df['TSH'].quantile(0.25)
Q3_TSH = df['TSH'].quantile(0.75)
IQR_TSH = Q3_TSH - Q1_TSH
lower_TSH = df['TSH'] < (Q1_TSH - 3 * IQR_TSH)
upper_TSH = df['TSH'] > (Q3_TSH + 3 * IQR_TSH)
print('TSH:', 'lower outliers -', sum(lower_TSH), ' | upper outliers -', sum(upper_TSH))

# T3
Q1_T3 = df['T3'].quantile(0.25)
Q3_T3 = df['T3'].quantile(0.75)
IQR_T3 = Q3_T3 - Q1_T3
lower_T3 = df['T3'] < (Q1_T3 - 3 * IQR_T3)
upper_T3 = df['T3'] > (Q3_T3 + 3 * IQR_T3)
print('T3:', 'lower outliers -', sum(lower_T3), ' | upper outliers -', sum(upper_T3))

# TT4
Q1_TT4 = df['TT4'].quantile(0.25)
Q3_TT4 = df['TT4'].quantile(0.75)
IQR_TT4 = Q3_TT4 - Q1_TT4
lower_TT4 = df['TT4'] < (Q1_TT4 - 3 * IQR_TT4)
upper_TT4 = df['TT4'] > (Q3_TT4 + 3 * IQR_TT4)
print('TT4:', 'lower outliers -', sum(lower_TT4), ' | upper outliers -', sum(upper_TT4))

# T4U
Q1_T4U = df['T4U'].quantile(0.25)
Q3_T4U = df['T4U'].quantile(0.75)
IQR_T4U = Q3_T4U - Q1_T4U
lower_T4U = df['T4U'] < (Q1_T4U - 3 * IQR_T4U)
upper_T4U = df['T4U'] > (Q3_T4U + 3 * IQR_T4U)
print('T4U:', 'lower outliers -', sum(lower_T4U), ' | upper outliers -', sum(upper_T4U))

# FTI
Q1_FTI = df['FTI'].quantile(0.25)
Q3_FTI = df['FTI'].quantile(0.75)
IQR_FTI = Q3_FTI - Q1_FTI
lower_FTI = df['FTI'] < (Q1_FTI - 3 * IQR_FTI)
upper_FTI = df['FTI'] > (Q3_FTI + 3 * IQR_FTI)
print('FTI:', 'lower outliers -', sum(lower_FTI), ' | upper outliers -', sum(upper_FTI))

TSH: lower outliers - 0 | upper outliers - 456
T3: lower outliers - 0 | upper outliers - 84
TT4: lower outliers - 0 | upper outliers - 55
T4U: lower outliers - 7 | upper outliers - 22
FTI: lower outliers - 0 | upper outliers - 83

```



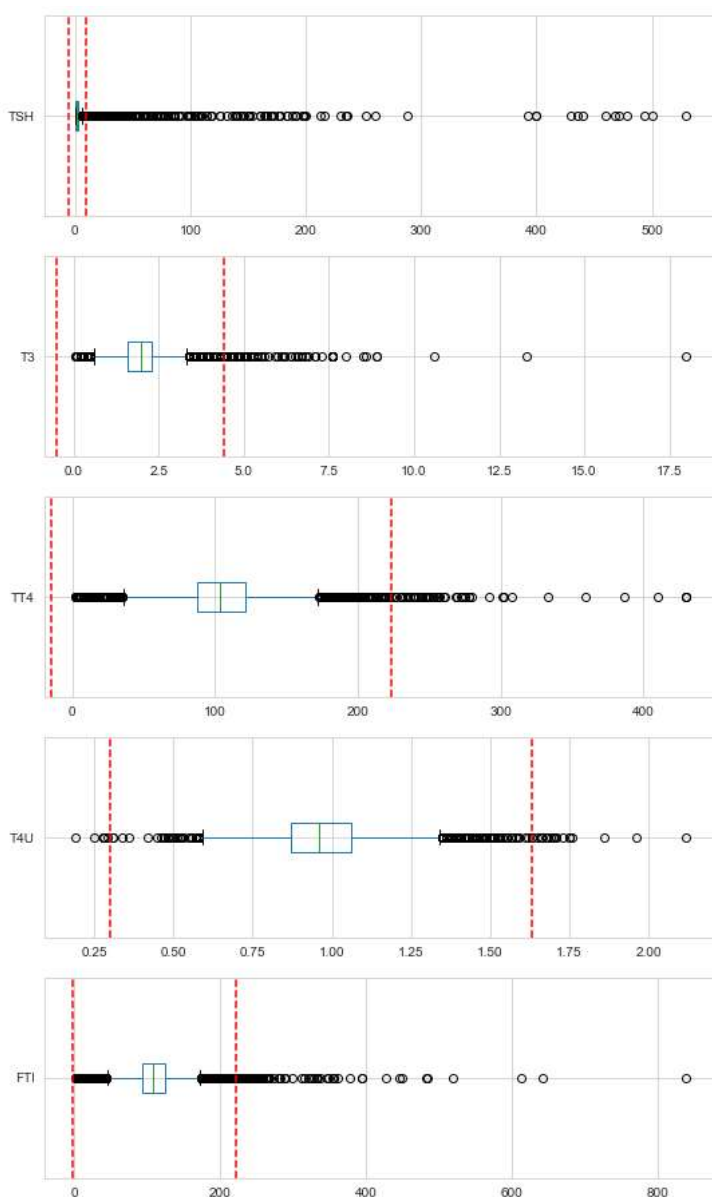
```

In [23]: # preparing boxplots
fig, axs= plt.subplots(nrows = 5, figsize=(9,16))

# TSH
df.boxplot(column='TSH', ax=axs[0], vert = False)
axs[0].axvline(x=(Q1_TSH - 3*IQR_TSH), color='r', linestyle='--')
axs[0].axvline(x=(Q3_TSH + 3*IQR_TSH), color='r', linestyle='--')
# T3
df.boxplot(column='T3', ax=axs[1], vert = False)
axs[1].axvline(x=(Q1_T3 - 3*IQR_T3), color='r', linestyle='--')
axs[1].axvline(x=(Q3_T3 + 3*IQR_T3), color='r', linestyle='--')
# TT4
df.boxplot(column='TT4', ax=axs[2], vert = False)
axs[2].axvline(x=(Q1_TT4 - 3*IQR_TT4), color='r', linestyle='--')
axs[2].axvline(x=(Q3_TT4 + 3*IQR_TT4), color='r', linestyle='--')
# T4U
df.boxplot(column='T4U', ax=axs[3], vert = False)
axs[3].axvline(x=(Q1_T4U - 3*IQR_T4U), color='r', linestyle='--')
axs[3].axvline(x=(Q3_T4U + 3*IQR_T4U), color='r', linestyle='--')
# FTI
df.boxplot(column='FTI', ax=axs[4], vert = False)
axs[4].axvline(x=(Q1_FTI - 3*IQR_FTI), color='r', linestyle='--')
axs[4].axvline(x=(Q3_FTI + 3*IQR_FTI), color='r', linestyle='--')

```

Out[23]: <matplotlib.lines.Line2D at 0x1e31d169fd0>



It seems that we have many severe outliers present. However, given our research about these values, this type of variance is normal within this context and is to be expected, especially when dealing with persons who will have alterations in these values given their medical conditions.

**Observations:**

In [24]: *#--> It doesnt seem Like a coincidence that most missing values present are from blood tests*  
*#--> We need to investigate this further in order to decide the best approach to handling them*

**calculating missingess of entire DF**

In [25]:

```
missingness = df.isnull().sum().sum() / df.count().sum()
print('Overall Missingness of thyroidDF is: {:.2f}%'.format(missingness * 100))

# Create table for missing data analysis
def missing_table(df):
    total = df.isnull().sum().sort_values(ascending=False)
    percent = (df.isnull().sum()/df.isnull().count()).sort_values(ascending=False)
    missing_data = pd.concat([total, percent], axis=1, keys=['Total', 'Percent'])
    return missing_data

# Analyze missing data
missing_table(df).head(10)
```

Overall Missingness of thyroidDF is: 7.54%

Out[25]:

	Total	Percent
TBG	7287	0.965677
T3	2209	0.292738
TSH	722	0.095680
T4U	676	0.089584
FTI	669	0.088656
TT4	354	0.046912
sex	250	0.033130
age	4	0.000530
goitre	0	0.000000
psych	0	0.000000

**Missing Table Summary**

In [26]: *# thyroidDF['TBG'] - 96.56% missing*  
*# The 'TBG' attribute is almost entirely missing from the dataset. This column will be removed at once!*

*# thyroidDF['age'] - 0.045% missing*  
*# We will also go ahead and drop these 4 observations from the dataset. All 4 of these observations belong to observations with*

*# thyroidDF['Sex'] - 3.37% missing*  
*# There are a total of 300 observations where 'sex' is null. In an attempt to preserve some of these values,*  
*#we check how many of these observations also have 'pregnant' == True.*  
*#There are 3 such observations. Assuming the 'pregnancy' attribute is correct for these observations, we can confidently say these*


In [27]: *# dropping 'TBG' attribute from dataset*  
df.drop(['TBG'], axis=1, inplace=True)

*# dropping 4 observations with abnormal 'age' from dataset*  
df.dropna(subset=['age'], inplace=True)

*# changing sex of observations with ('pregnant' == True) & ('sex' == null) to Female*  
df['sex'] = np.where((df.sex.isnull()) & (df.pregnant == 't'), 'F', df.sex)

Now lets take a look at missing values per row. If we are moving forward with imputation, we dont want to keep rows that have too many missing values (especially since most missing values in the dataset are from the most important attributes).

```
In [28]: # count missing values per row
df['n_missing'] = df.isnull().sum(axis=1)
sns.histplot(df, x='n_missing', binwidth=0.5, hue='target');
```



It looks like after removing TBG from the dataset, most of the observations in our dataset have zero missing values. A lot have one missing value and some have two. a minority have 3 or more. Before we do any imputation we will remove the observations that are missing data for more than 2 columns.

```
In [29]: # calculating missingness of entire DF
missingness = df.isnull().sum().sum() / df.count().sum()
print('Overall Missingness of thyroidDF is: {:.2f}%'.format(missingness * 100))

# remove rows with 3 or more missing values
df.drop(df.index[df['n_missing'] > 2], inplace=True)
print
missing_table(df).head(10)
```

Overall Missingness of thyroidDF is: 2.89%

Out[29]:

	Total	Percent
<b>T3</b>	1910	0.267432
<b>TSH</b>	406	0.056847
<b>T4U</b>	290	0.040605
<b>FTI</b>	283	0.039625
<b>sex</b>	208	0.029123
<b>TT4</b>	6	0.000840
<b>age</b>	0	0.000000
<b>goitre</b>	0	0.000000
<b>target</b>	0	0.000000
<b>psych</b>	0	0.000000

```
In [30]: # replacing boolean strings with binary 0 and 1
df.replace('f', 0, inplace=True)
df.replace('t', 1, inplace=True)

# replacing sex with binary 0 and 1
df.replace('M', 0, inplace=True) # male mapped to 0
df.replace('F', 1, inplace=True) # female mapped to 1

thyroid_df = df.replace(np.nan, 0)

# re-mapping target vaues to diagnostic groups
# *** I get different final results by ordering these classes differently ***
diagnoses = {'negative': 0,
             'hypothyroid': 1,
             'hyperthyroid': 2}
thyroid_df['target'] = thyroid_df['target'].map(diagnoses)
```

In [31]: thyroid\_df

Out[31]:

	age	sex	on_thyroxine	query_on_thyroxine	on_antithyroid_meds	sick	pregnant	thyroid_surgery	I131_treatment	query_hypothyroid	...	tumor	hypopit
1	29.0	1.0	0	0	0	0	0	0	0	0	...	0	0
7	28.0	1.0	0	0	0	0	0	0	0	0	...	0	0
8	28.0	1.0	0	0	0	0	0	0	0	0	...	0	0
9	28.0	1.0	0	0	0	0	0	0	0	0	...	0	0
10	54.0	1.0	0	0	0	0	0	0	0	0	...	0	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9166	70.0	1.0	0	0	0	0	0	0	0	0	...	0	0
9167	56.0	0.0	0	0	0	0	0	0	0	0	...	0	0
9168	22.0	0.0	0	0	0	0	0	0	0	0	...	0	0
9170	47.0	1.0	0	0	0	0	0	0	0	0	...	0	0
9171	31.0	0.0	0	0	0	0	0	0	0	0	...	1	0

7142 rows × 23 columns

### Traning Data:

```
In [32]: from sklearn.model_selection import train_test_split
X = thyroid_df.drop('target', axis=1).copy()
y = thyroid_df['target'].copy()

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42, stratify=y)
```

### Model Creation:

```
In [33]: # Import the model we are using
from sklearn.ensemble import RandomForestClassifier
# Instantiate model with 100 decision trees
rc_model = RandomForestClassifier(n_estimators=100, criterion='entropy')

rc_model.fit(X_train, y_train)

rc_pred = rc_model.predict(X_test)
```

### classification\_report:

```
In [34]: from sklearn.metrics import confusion_matrix, accuracy_score, classification_report

print('classification_report:')
print(classification_report(y_test, rc_pred))

acc = accuracy_score(y_test, rc_pred)
print('accuracy of the model :', acc)
```

```
classification_report:
              precision    recall  f1-score   support

     0       0.99      0.99      0.99       1597
     1       0.95      1.00      0.97        145
     2       0.89      0.77      0.83         44

   accuracy                   0.99       1786
  macro avg       0.95      0.92      0.93       1786
 weighted avg       0.99      0.99      0.99       1786
```

accuracy of the model : 0.9876819708846585

In [ ]: